



# REAL ESTATE PRICING PREDICTIONS

Elio Aybar, Cristal Garcia, Sunny Li,  
and Matt Norgren



**GOAL:** PREDICT MEDIAN HOUSE PRICES GIVEN CONVENTIONAL ATTRIBUTES AND PRESENCE OF A TARGET AND/OR STARBUCKS.



**VALUE:** TESTING NON-CONVENTIONAL FEATURES TO SEE IF WORTH EXPLORING IN THE FUTURE.



**HYPOTHESIS:** DIFFERENT TYPES OF NEIGHBORHOODS MAY ATTRACT SPECIFIC BUSINESSES.

# PROBLEM STATEMENT





# PROJECT DEFINITION OF VARIABLES

- Zillow Real Estate: 4.3M X 75
- Starbucks location 25K X 7
- Target Locations: 1.8K X 47
- USA Public Schools data: 47K X 33
- Food Desert & Economic Data: 73K X 148
- Hospital Ratings: 4.8K X 28



# FEATURES & TARGET VARIABLE

## Zip code (Key)



## Zillow

- NumberOfBedrooms
- Dwelling Type



## School

- LEVEL\_: Grade level
- Class\_Teacher\_RATIO:



## Food Desert/Economic Data

- Urban Indicator
- Occupied Unit Housing Count (2010 Census)
- Low Income Indicator
- Median Family Income
- Food Desert Indicator



## Hospital

- Hospital Rating
- Emergency Services Flag



## Starbucks Locations



## Target Locations



## Pricing



# LIMITATIONS AND DELIMITATIONS OF DATA

---

01

Hospital:  
hospital deserts  
(distance between  
hospital and  
home)

02

Public School  
Data: Creating  
Student-to-  
Teacher Ratio

03

Food Desert  
Data: converting  
CensusTract  
number to  
Zipcode

04

Starbucks and  
Target



# LIMITATIONS AND DELIMITATIONS OF DATA PART 2



Granularity



Time

(1996 – 2017)



Features

Sales Data

Zillow Estimates



Complexity of combination

19K Unique Zip codes



Aggregation of Features

Number of Bedrooms

Dwelling Type



# IS IT ALL HERE?



Scarcity of price



Duplicative feature types



Scaling & Aggregation



Consolidated variable creation

Function creation  
Factorization



Melting and Reshaping

```
p_df['MedianListingPricePerSqft_1Bedroom'] = zip_
p_df['MedianListingPricePerSqft_2Bedroom'] = zip_
p_df['MedianListingPricePerSqft_3Bedroom'] = zip_
p_df['MedianListingPricePerSqft_4Bedroom'] = zip_
p_df['MedianListingPricePerSqft_5BedroomOrMore'] =
```

```
p_df['MedianListingPrice_1Bedroom'] = zip_df['Medi
p_df['MedianListingPrice_2Bedroom'] = zip_df['Medi
p_df['MedianListingPrice_3Bedroom'] = zip_df['Medi
p_df['MedianListingPrice_4Bedroom'] = zip_df['Medi
p_df['MedianListingPrice_5BedroomOrMore'] = zip_c
```

```
df['MedianRentalPricePerSqft_1Bedroom'] = zip_c
df['MedianRentalPricePerSqft_2Bedroom'] = zip_c
df['MedianRentalPricePerSqft_3Bedroom'] = zip_c
df['MedianRentalPricePerSqft_4Bedroom'] = zip_c
df['MedianRentalPricePerSqft_5BedroomOrMore'] =
```

```
ed: [ ] [ ] [ ] [ ] [ ] 1Bedroom'] = zip_df['Medi
ed: [ ] [ ] [ ] [ ] [ ] 2Bedroom'] = zip_df['Medi
anRentalPrice_3Bedroom'] = zip_df['Medi
yRentalPrice_4Bedroom'] = zip_df['Medi
entalPrice_5BedroomOrMore'] = zip_df
```

```
ingPrice_AllHomes'] = zip_df['Medi
gPrice_CondoCoop'] = zip_df['Me
rice_DuplexTriplex'] = zip_df
e_SingleFamilyResidence']
```





# DATA EXPLORATION

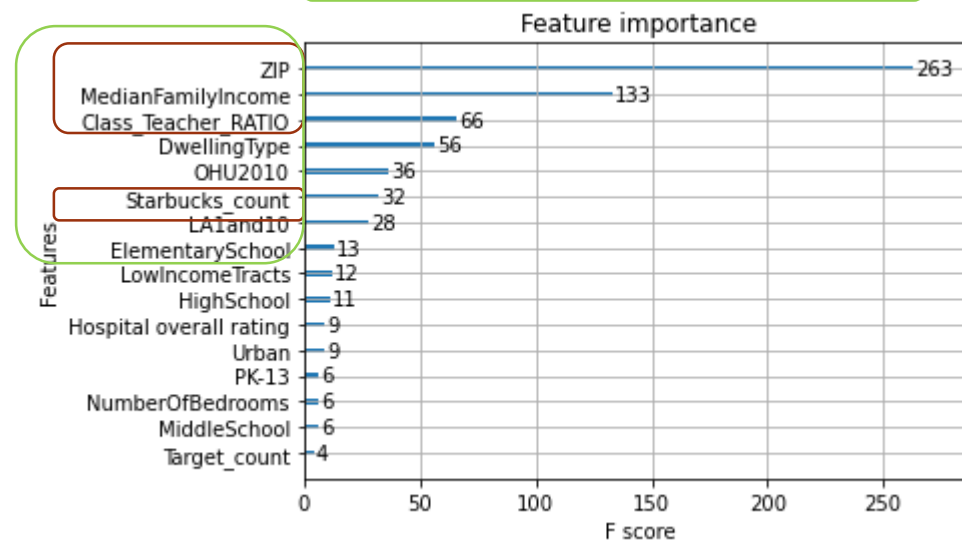


## Featurization:



- 90%-time spent feature engineering (5.5 weeks)
- Individually worked on datasets separately:
  - Identified and removed features with high collinearity
  - Captured summary statistics
- Consolidated features and visualized relationship





# HEATMAP, PCA & FEATURE IMPORTANCE





# TARGET VARIABLE EXPLORATION

```
Balance data using SMOTE

# pip install --upgrade scikit-learn

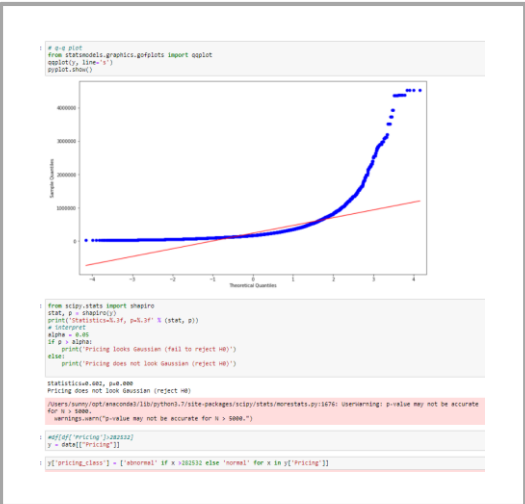
# pip install imblearn
from imblearn.over_sampling import SMOTE
sm = SMOTE(sampling_strategy='auto')

X_sm, y_sm = sm.fit_sample(X, y['pricing_class'])
X_sm.shape, y_sm.shape

((92576, 18), (92576,))

from collections import Counter
print('Original dataset shape {}'.format(Counter(y_sm)))
print('Resampled dataset shape {}'.format(Counter(y_sm)))

Original dataset shape Counter({'normal': 46288, 'abnormal': 46288})
Resampled dataset shape Counter({'normal': 46288, 'abnormal': 46288})
```



# METHODOLOGY

- Models
  - Linear Regression
  - Decision Tree
  - Random Forest
  - Support Vector Regression
  - Gradient Boosting
  - KNN for Regression
  - XGBoost





1

## Validation Methodology

- Compare Desired Errors with Training and Cross Validation errors

2

## Descriptions

- Choose top 3 models

3

## Model Selection Justification

- Best performing RMSE

# FRAMEWORK





[This Photo](#) by Unknown Author is licensed under [CC BY-SA](#)

# RESULTS

# MODELING RESULTS

- Linear Regression
  - Desired Error: 279471.46
  - Training Error: 0.78
  - Cross Validation Error: 0.78
- Support Vector Regression (SVR)
  - RMSE score: 0.65
- Decision Tree
  - Hyperparameter Tuning: max depth = 4, min\_samples\_leaf = 0.1
  - Desired Error: 210381.81
  - Training Error: 194304.60
  - Cross Validation Error: 195936.04
- Random Forest
  - Hyperparameter Tuning: max depth = 4, min\_samples\_leaf = 0.1
  - Desired Error: 204785.74
  - Training Error: 201373.10
- Gradient Boosting
  - RMSE: 131471.88186
- KNN for Regression
  - N\_neighbors = 2, RMSE value for k = 2 of 325245.033
- XGBoost
  - RMSE = 108618





# MODELING RESULTS DIFFERENT FEATURES

- Linear Regression
  - Desired Error: 294214.72
  - Training Error: 0.84
  - Cross Validation Error: 0.84
- Support Vector Regression (SVR)
  - RMSE score: 0.79
- Decision Tree
  - Hyperparameter Tuning: max depth = 4, min\_samples\_leaf = 0.1
  - Desired Error: 204245.67
  - Training Error: 200400.99
  - Cross Validation Error: 200487.57
- Random Forest
  - Hyperparameter Tuning: max depth = 4, min\_samples\_leaf = 0.1
  - Desired Error: 208813.27
  - Training Error: 205483.11
- Gradient Boosting
  - RMSE: 173276.42
- KNN for Regression
  - N\_neighbors = 2, RMSE value for k = 2 of 327642.71
- XGBoost
  - RMSE = 112413



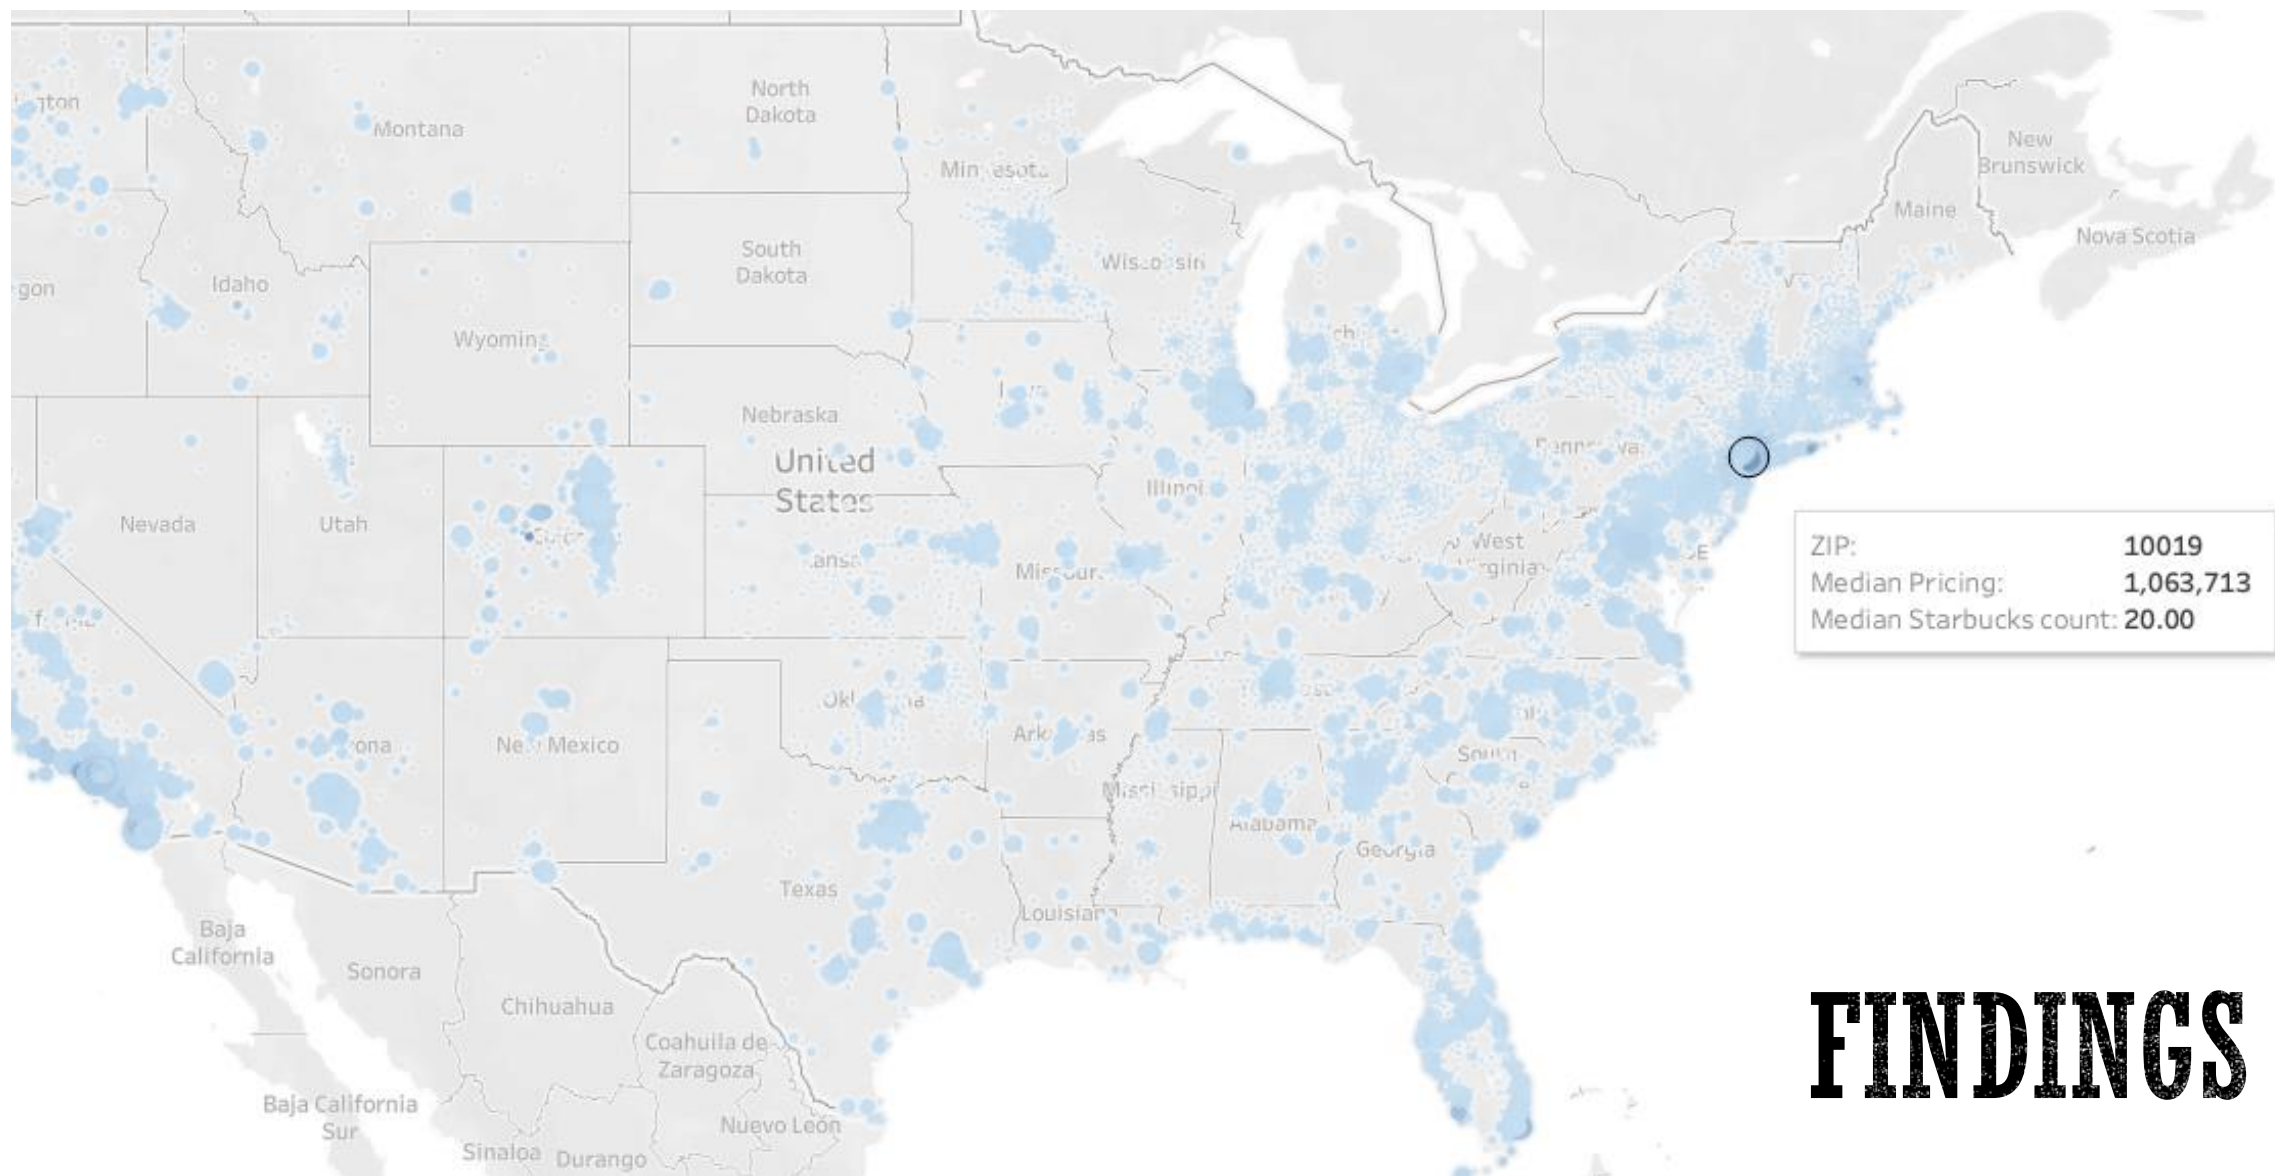


## eXtreme Gradient Boosted trees

Most powerful machine learning algorithm up until today

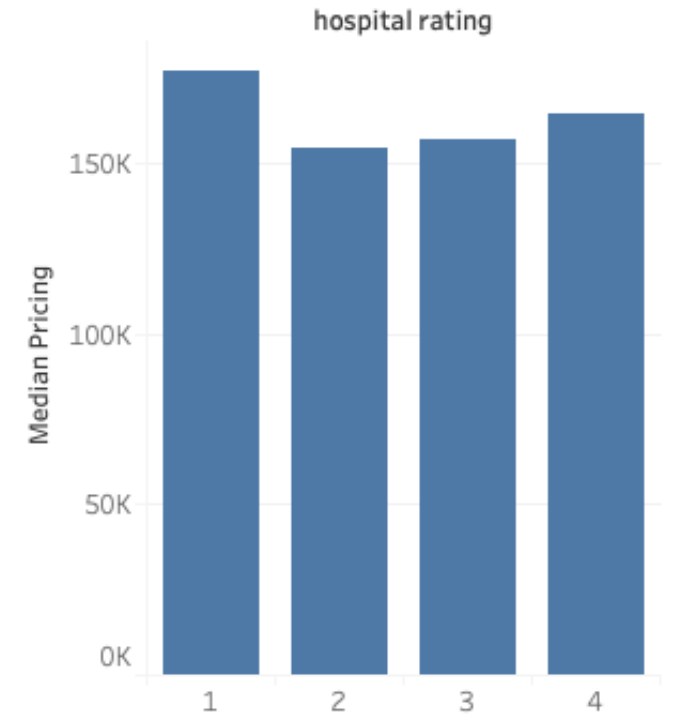

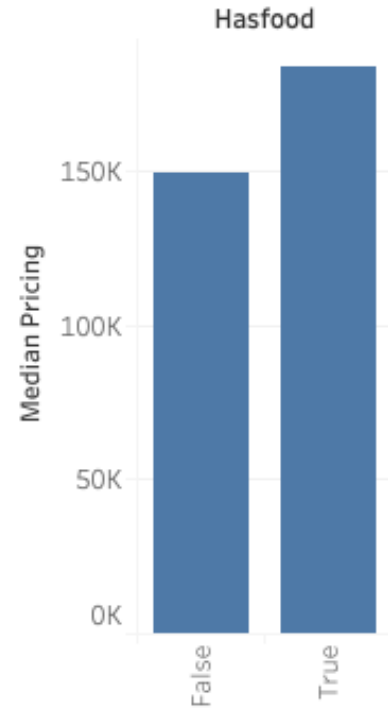
### Features:

- Regularized boosting ( prevent overfitting)
- Can handle missing values
- Parallel processing
- Can cross-validation at each iteration




# FINDINGS

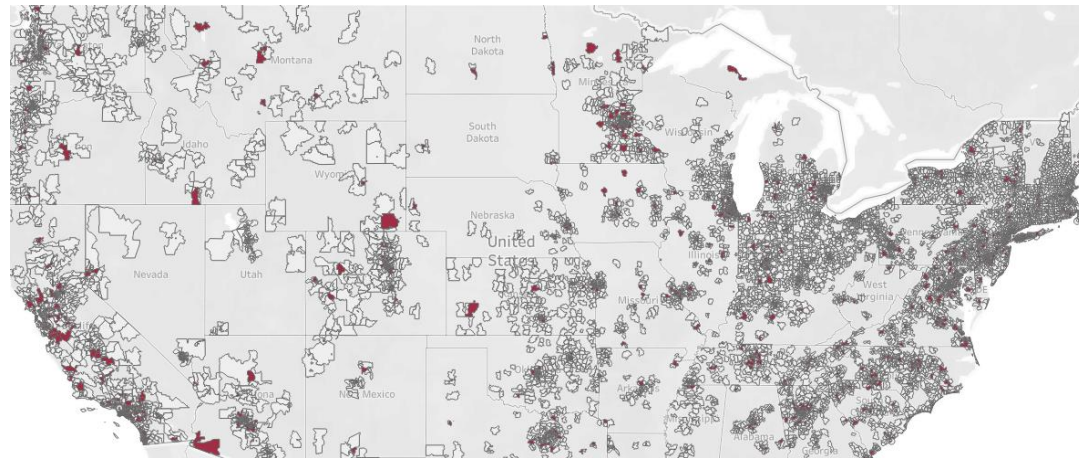




# CONVENTIONAL FEATURES



## Target Density



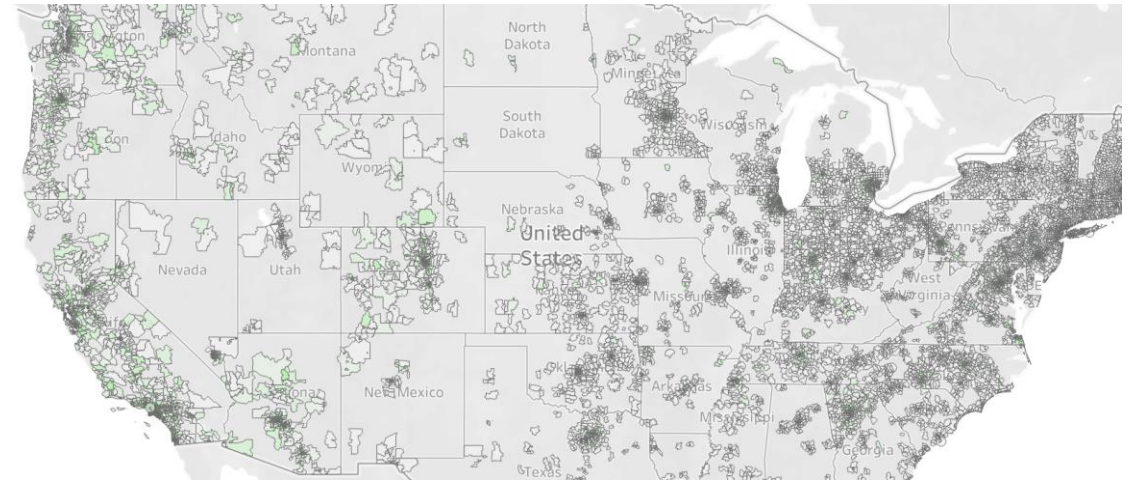
MEDIAN(Target count)




0.000

1.000

## Starbucks Density




MEDIAN(Starbucks count)



0.00

21.00

# STARBUCKS INFLUENCE








Non-conventional features may be economic indicators.



We spent 80% of time on addressing a number of issues around data quality, standards, access.



We have plans for model optimization but didn't have time to achieve it.

## CONCLUSIONS



# FUTURE CONSIDERATIONS

- New dataset or extension of features.
- Number of bathrooms, rent cost, Square footage, etc.
- Dollar trees stores, Dunkin donuts, pharmacies, supermarkets, etc
- Complex imputation methods
  - Focus on the hospital specifically, calculate the radii of hospitals, then use radii to weigh the hospital rating for each zip code.
- Robust test
  - Replacing zip code with county/city
- Include time series data to the analysis





# QUESTIONS?

Thanks for Your attention

